



# Reinforcement Learning

PhD course @ “Dottorato di ingegneria industriale e dell’informazione”  
Trieste, 2024



Simone Silveti  
Research & Development

# Who am I?



**Simone Silvetti** (silvetti@esteco.com)

- Studied mathematics in Rome
- Phd in Computer Science @ Udine
- Currently working in ESTECO

application of quantitative  
formal methods and  
machine learning techniques  
to Verification and  
Model-based Testing of  
Complex Systems



# Who am I?



**Simone Silvetti** (silvetti@esteco.com)

- Studied mathematics in Rome
- Phd in Computer Science @ Udine
- Currently working in **ESTECO**

application of quantitative formal methods and machine learning techniques to Verification and Model-based Testing of Complex Systems

Numerical  
Methods Group

multi-objective optimization algorithms, machine learning, object-oriented programming

# Who am I?



**Simone Silvetti** (silvetti@esteco.com)

- Studied mathematics in Rome
- Phd in Computer Science @ Udine
- Currently working in **ESTECO**

application of quantitative formal methods and machine learning techniques to Verification and Model-based Testing of Complex Systems

Numerical  
Methods Group

multi-objective optimization algorithms, machine learning, object-oriented programming

Research and  
Development

process mining, research projects related to technology and domains useful for ESTECO products

# Who am I?



**Simone Silvetti** (silvetti@esteco.com)

- Studied mathematics in Rome
- Phd in Computer Science @ Udine
- Currently working in **ESTECO**

application of quantitative formal methods and machine learning techniques to Verification and Model-based Testing of Complex Systems

I worked on “Inverse Reinforcement Learning” applied to autonomous driving

Numerical  
Methods Group

multi-objective optimization algorithms, machine learning, object-oriented programming

Research and  
Development

process mining, research projects related to technology and domains useful for ESTECO products

# Who are you?



During your studies have you participated in courses of Reinforcement Learning? If yes, which topics have you covered?

13 responses

No

Only partially

I did not participate to any course.

I have never participated in a course about Reinforcement Learning.


I have never participated at any course of bayesian optimization

no

Foundations


Do you know Python? Numpy, Scipy?

13 responses




Will you follow the "Learning-based Controllers and the Reality Gap" course?

5 responses



# Reference

A book from Sutton et al.




Free available [here!](http://incompleteideas.net/book/the-book-2nd.html)

<http://incompleteideas.net/book/the-book-2nd.html>

# Reference

A book from Sutton et al.



Free available [here!](http://incompleteideas.net/book/the-book-2nd.html)

<http://incompleteideas.net/book/the-book-2nd.html>



**Google DeepMind**

@Google\_DeepMind · 482K subscribers · 186 videos

Artificial intelligence could be one of humanity's most useful inventions. Google DeepMind ... >

Subscribed

Home Videos Shorts Live Podcasts **Playlists** Community

Created playlists

Sort by



Inside Google DeepMind

View full playlist



Visualising AI

View full playlist




Scholarships | AI by you

View full playlist



Unfolded: Meet the scientists using AlphaFold

View full playlist



Life at DeepMind

View full playlist



The story of AlphaFold

View full playlist



Learning resources

View full playlist



Talks | AI for science

View full playlist



DeepMind: The Podcast - Season 2

View full podcast



DeepMind: The Podcast - Season 1

View full podcast



DeepMind x UCL | Deep Learning Lecture Series 2021

View full playlist



DeepMind x UCL | Deep Learning Lecture Series 2020

View full playlist




© ESTECO SpA



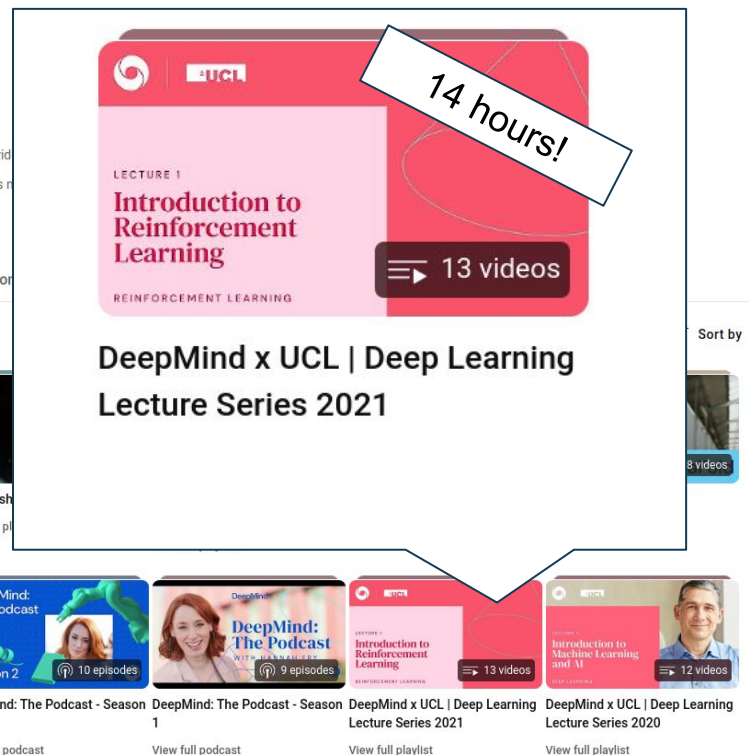
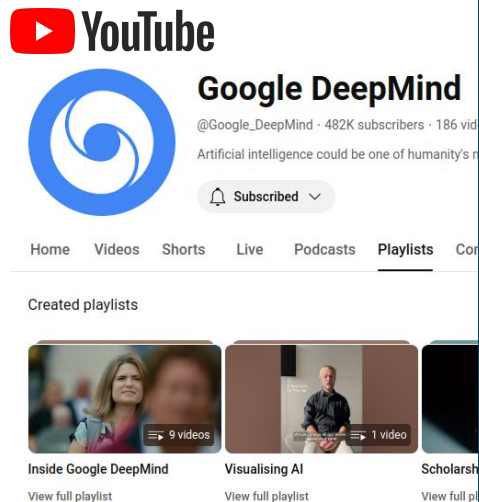
# Reference

A book from Sutton et al.




Free available [here!](http://incompleteideas.net/book/the-book-2nd.html)

<http://incompleteideas.net/book/the-book-2nd.html>



# Reference

A book from Sutton et al.



Free available [here!](http://incompleteideas.net/book/the-book-2nd.html)

<http://incompleteideas.net/book/the-book-2nd.html>




**ICTP Quantitative Life Sciences**

@ictpquantitativelifescienc9505 · 5.14K subscribers · 1K videos

More about this channel >

Subscribe

Home Videos Live Playlists Community



2020-2021 Reinforcement Learning (QLS-RL)

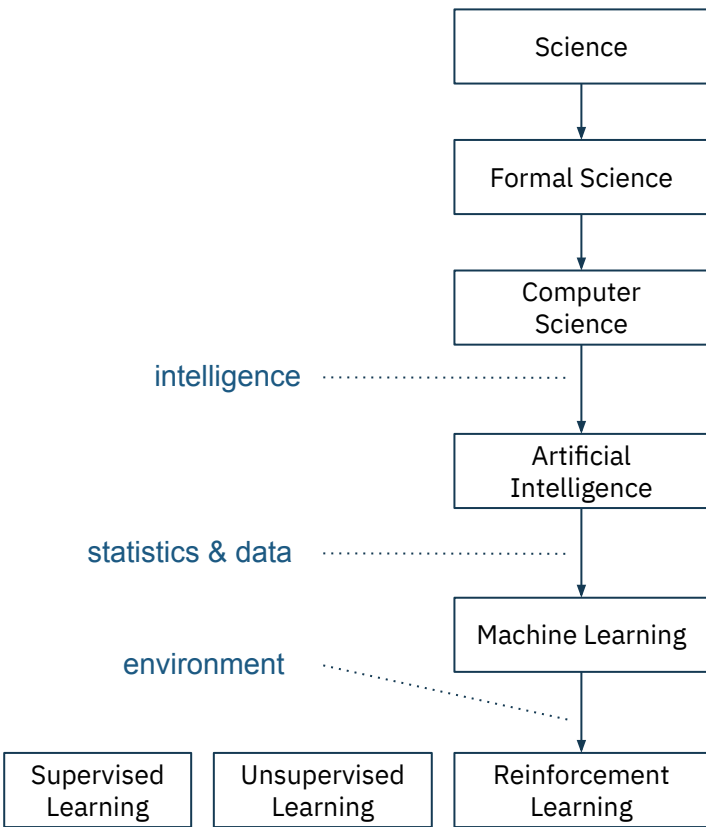
Prof. Antonio Celani

# Introduction

---

What is Reinforcement Learning?

# A map



the systematic study of physical and natural world through observation, experimentation, and the testing of theories against the evidence obtained

uses formal systems to generate knowledge

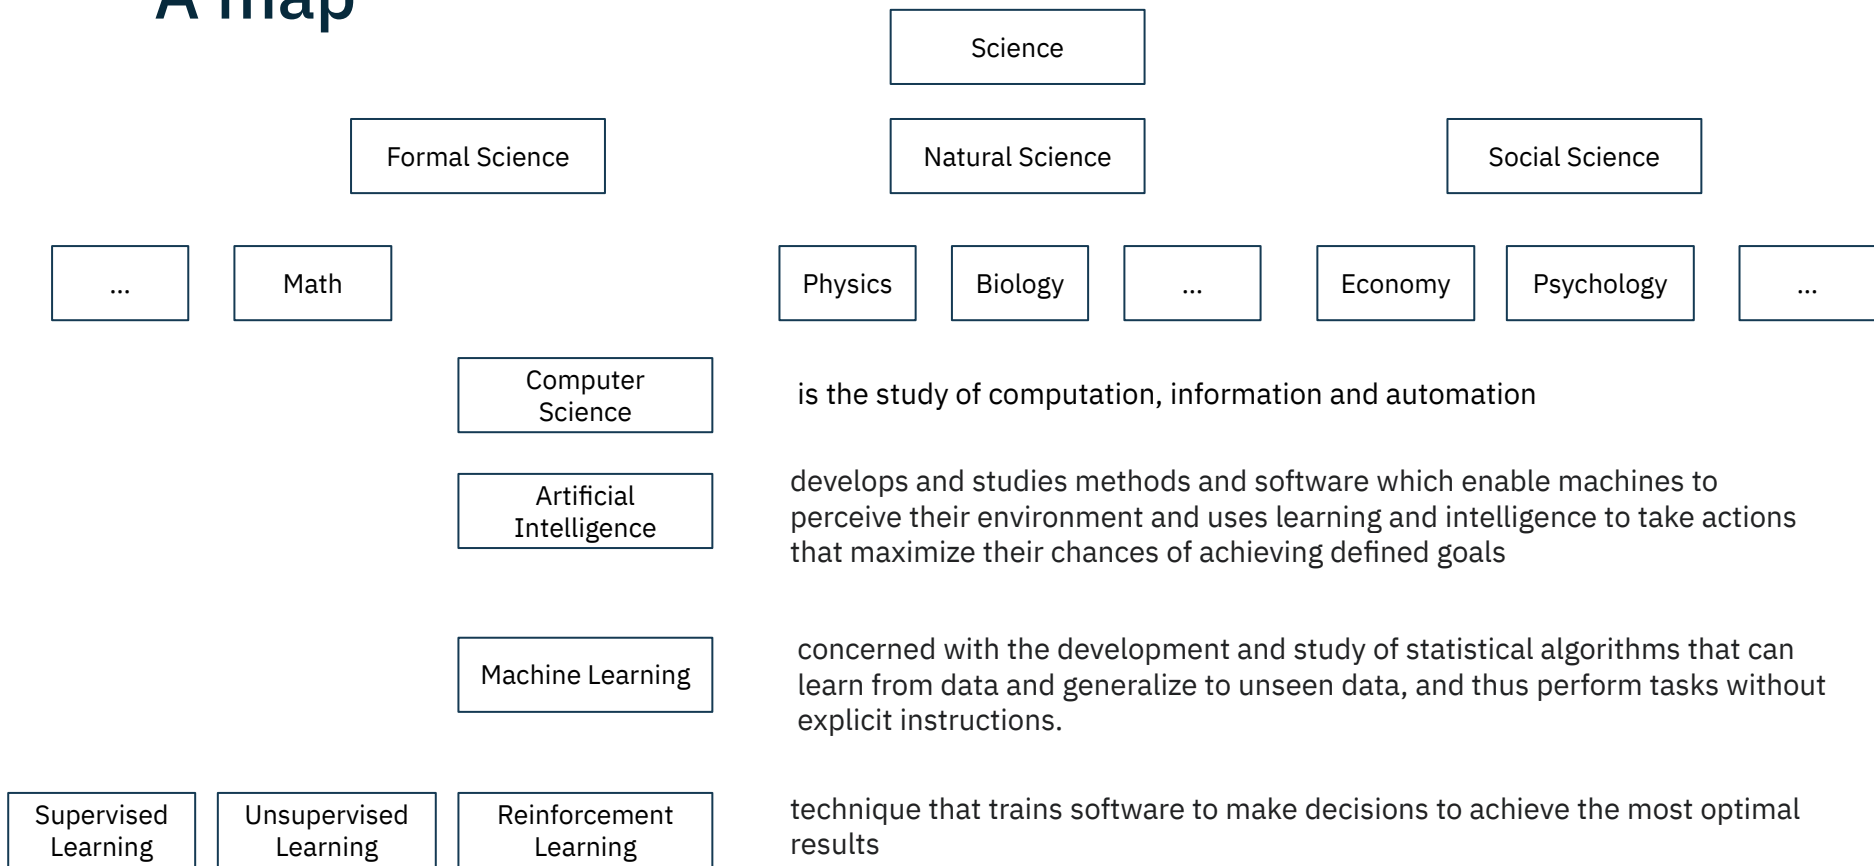
is the study of computation, information and automation

enabling machines to perceive their environment and uses learning and intelligence to take actions that maximize their chances of achieving defined goals

development and study of **statistical algorithms** that can learn from data and **generalize** to unseen data, and thus perform tasks without explicit instructions.

technique that trains software to make decisions to achieve the most optimal results

# A map



# A definition

Reinforcement  
Learning

technique that trains software to make decisions to achieve the most optimal results



# A definition

Reinforcement  
Learning

technique that trains software to make decisions to achieve the most optimal results



# A definition

Reinforcement  
Learning

technique that trains **agents** to make decisions to ~~achieve the most optimal results~~





# A definition

Reinforcement  
Learning

technique that trains **agents** to **map states into actions** to achieve the most optimal results



# A definition

Reinforcement  
Learning

technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**




# A definition

Reinforcement  
Learning

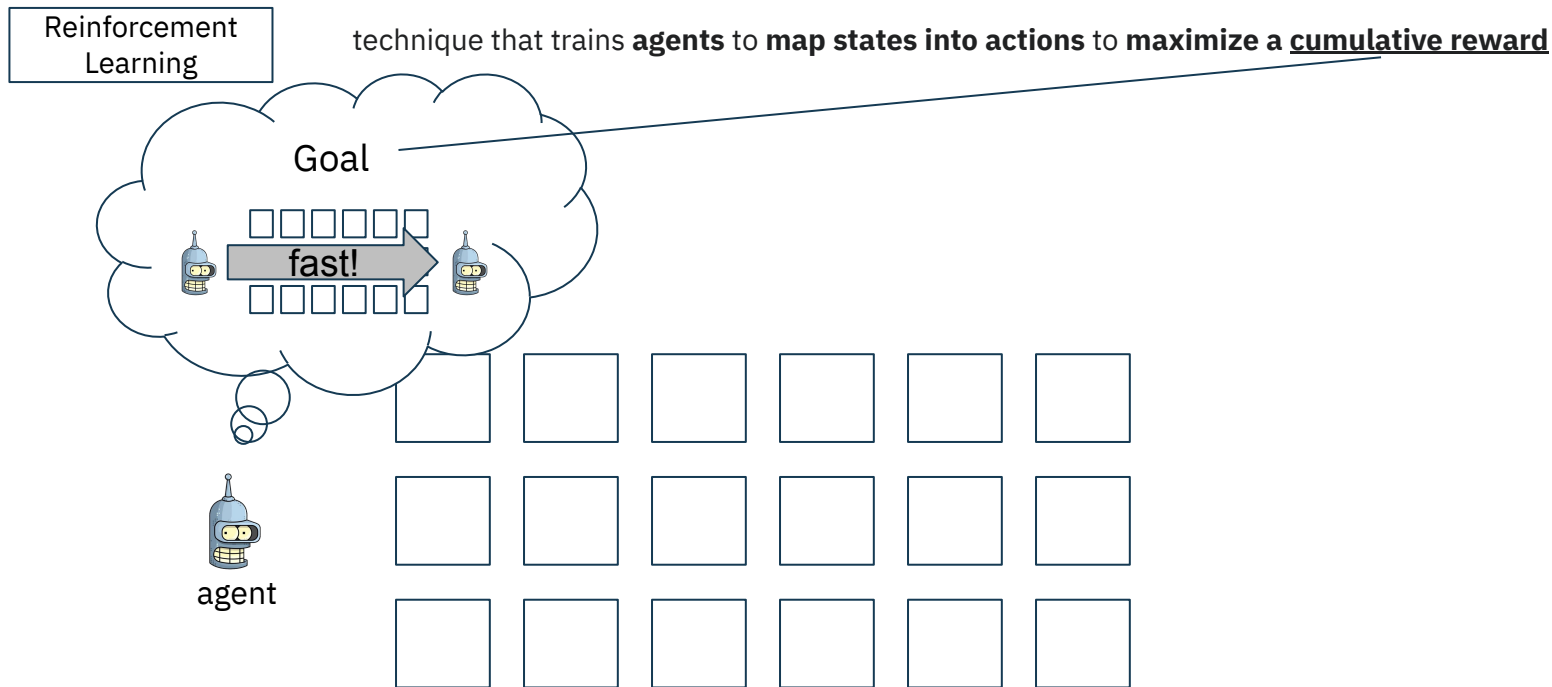
technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**



agent




# A definition



# A definition

Reinforcement  
Learning


technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**



# A definition

Reinforcement  
Learning


technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**



# A definition

Reinforcement  
Learning


technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**



# A definition

Reinforcement  
Learning

technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**






# A definition

Reinforcement  
Learning


technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**



# A definition

Reinforcement  
Learning


technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**



# A definition

Reinforcement  
Learning

technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**



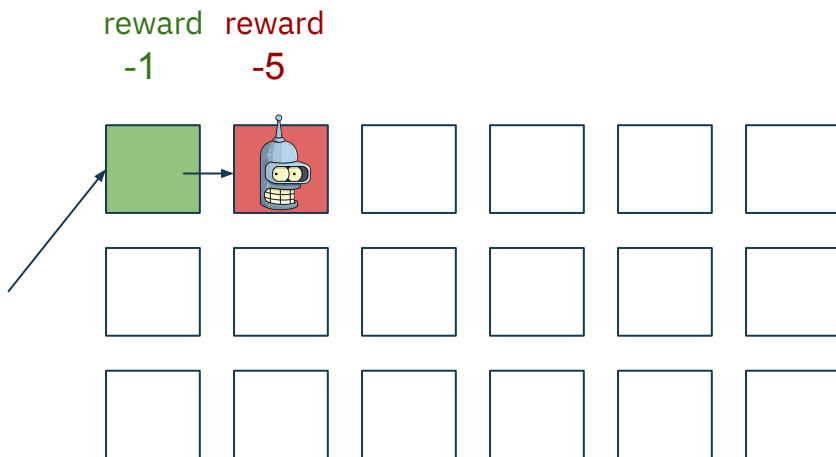
# A definition

Reinforcement  
Learning

technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**

Cumulative  
Reward


-6



# A definition

Reinforcement  
Learning

technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**



# A definition


Reinforcement  
Learning

technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**

Cumulative  
Reward

-5


reward  
-5



# A definition

Reinforcement  
Learning

technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**




# A definition

Reinforcement  
Learning

technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**

Cumulative  
Reward

-1






# A definition

Reinforcement  
Learning

technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**

Cumulative  
Reward

-6




# A definition

Reinforcement  
Learning

technique that trains **agents** to **map states into actions** to **maximize a cumulative reward**

Cumulative  
Reward  
94




# On reward

Goal and reward  
coherence

we want the agent goes as fast as possible from A to B. We need to choose an appropriate reward signal!

Cumulative  
Reward

94



# On reward


Goal and reward  
coherence

we want the agent goes as fast as possible from A to B. We need to choose an appropriate reward signal!

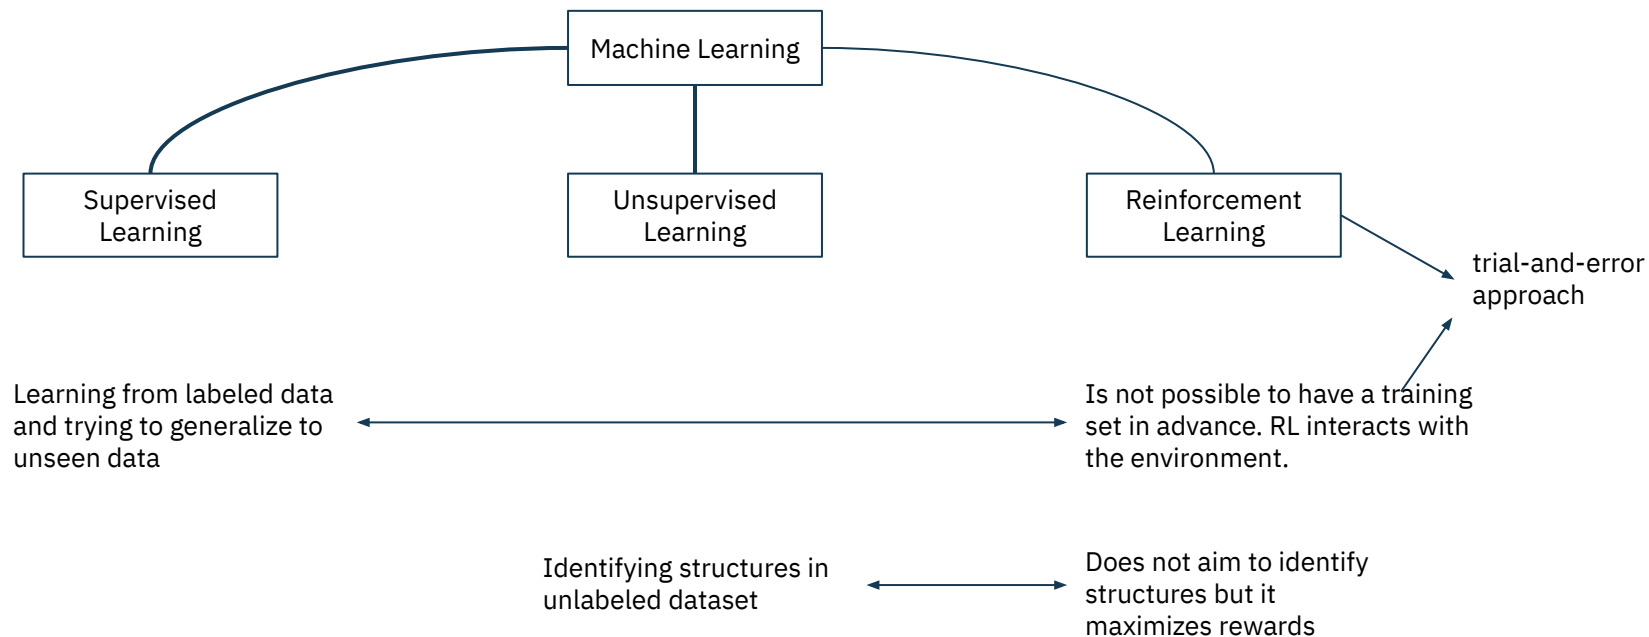
Cumulative  
Reward

94

Why negative values?



# A definition




# Elements of RL


---

Mathematical definition

# List of the ingredients



# Observability




environment




agent



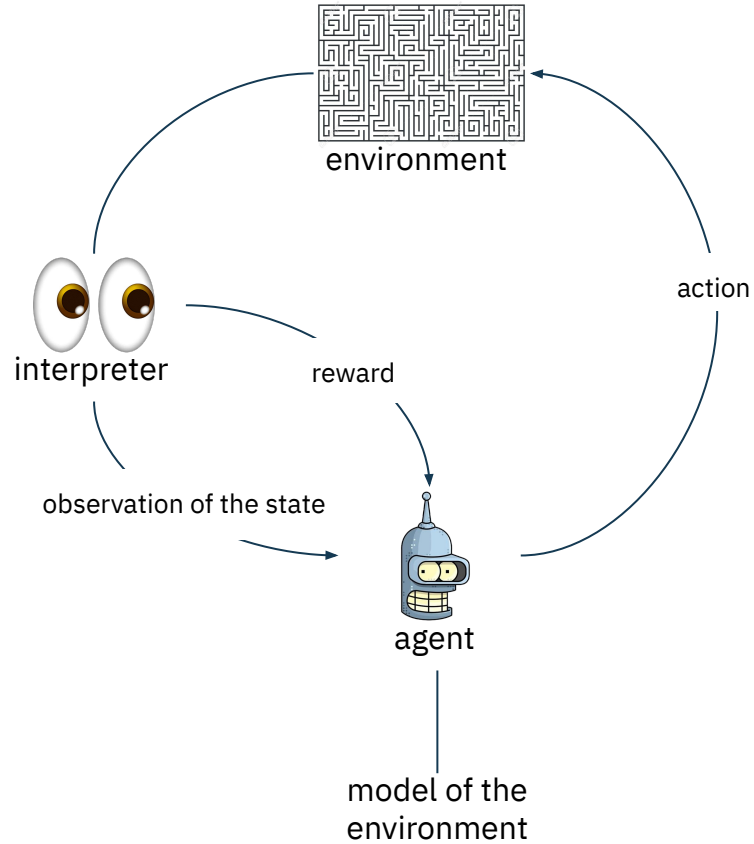
# Observability




# Observability



# Observability

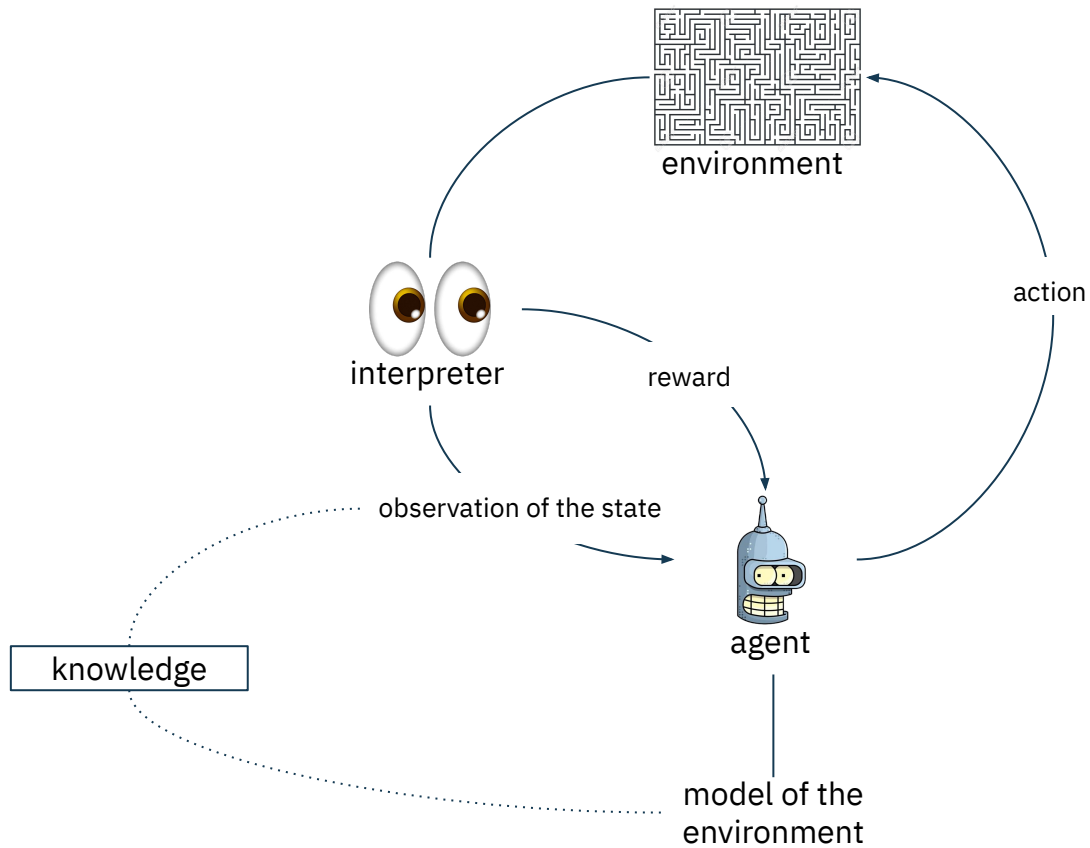


# Observability

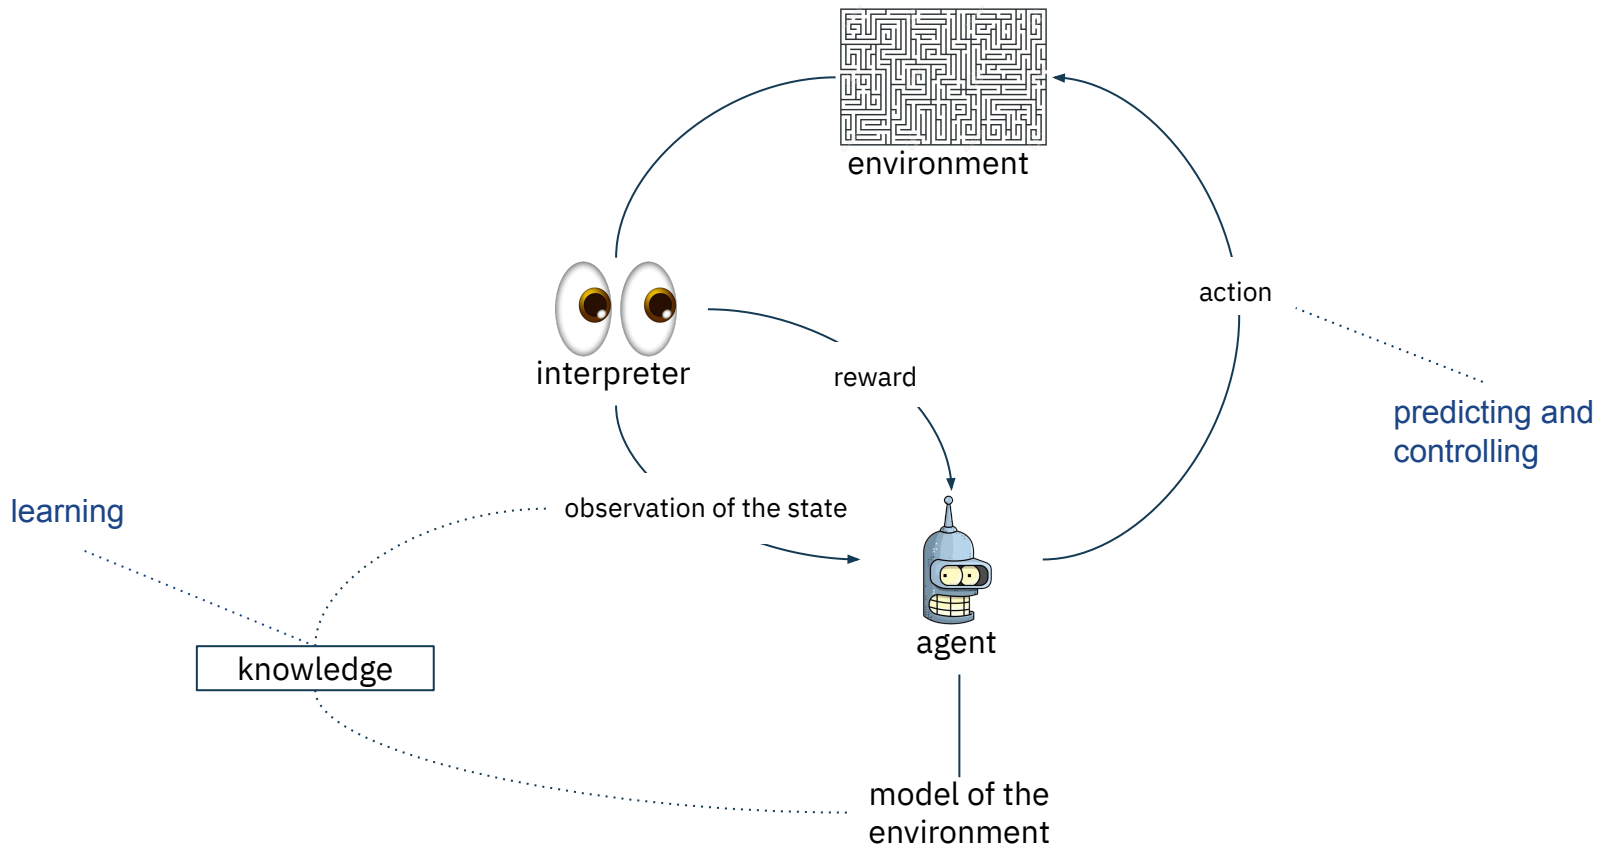


# Observability

RL = learning + prediction + controlling



# Observability



# Observability

RL = learning + prediction + controlling

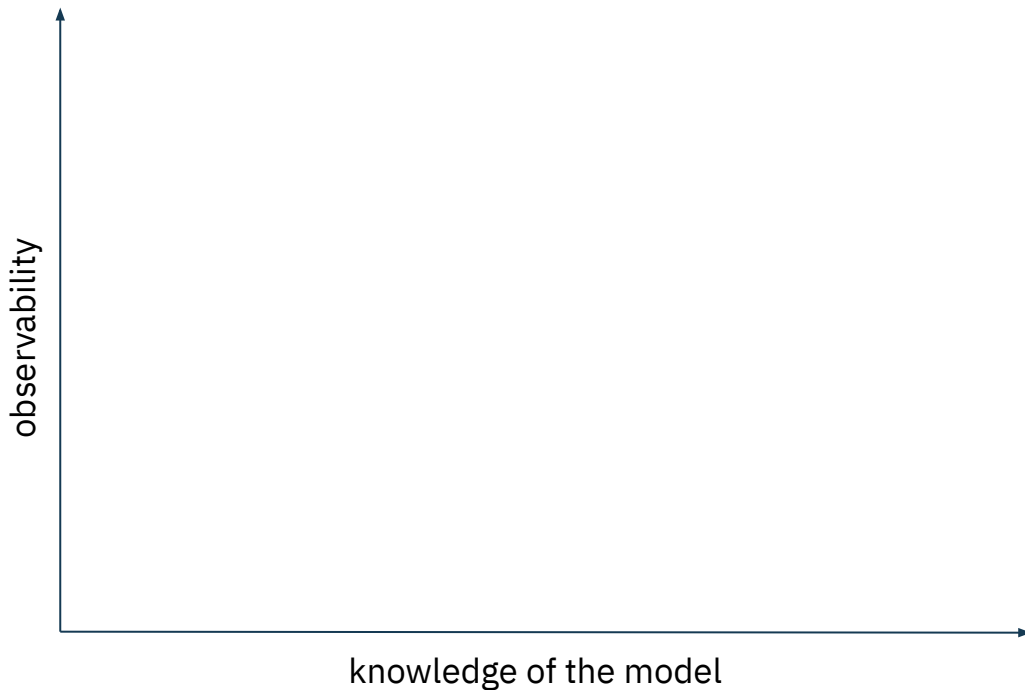
Building a model of  
the environment

Knowing the  
cumulative reward  
I'll get following a  
policy

Discovering the  
best action

# Knowledge of the environment

The two axes of knowledge



Empirical  
knowledge




Epistemic  
knowledge



# Knowledge of the environment

The two axes of knowledge




Markovian process  
only matter knowledge  
of the actual state

# Knowledge of the environment

The two axes of knowledge

Markovian process  
only matter knowledge  
of the actual state



Empirical  
knowledge



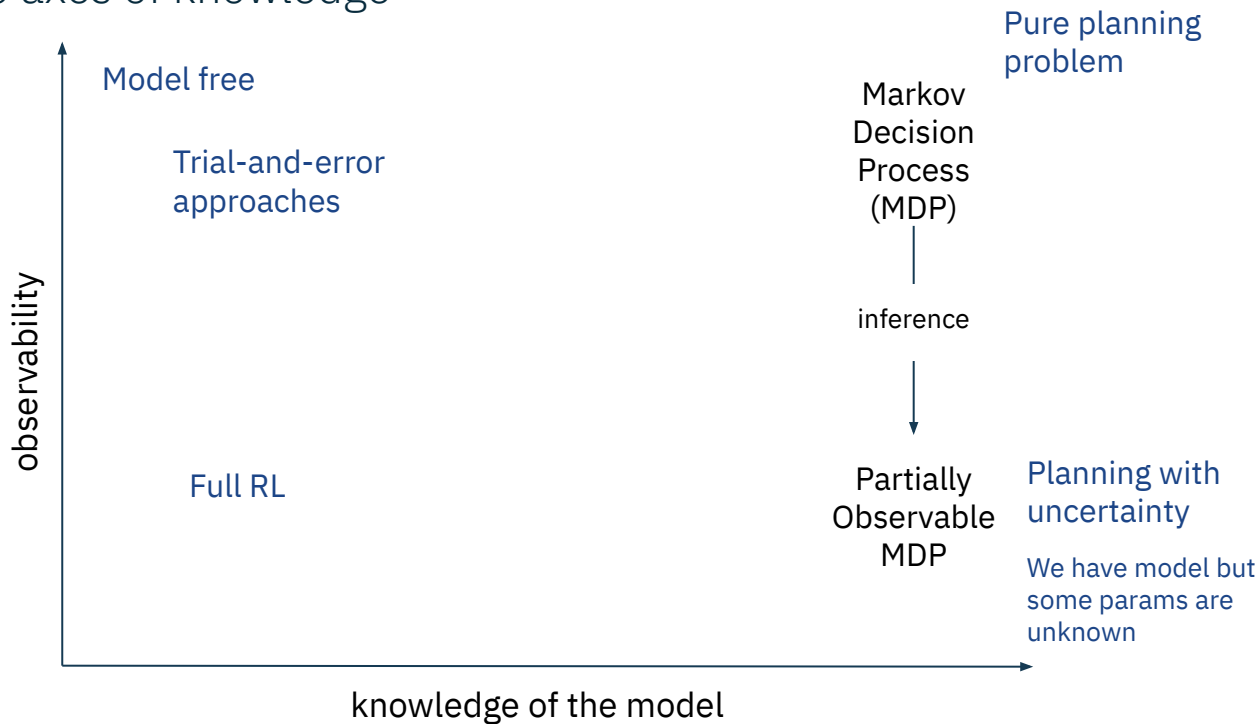
Epistemic  
knowledge



# Knowledge of the environment

The two axes of knowledge

Markovian process  
only matter knowledge  
of the actual state



Empirical  
knowledge




Epistemic  
knowledge



# Knowledge of the environment

The two axes of knowledge

Markovian process  
only matter knowledge  
of the actual state




Empirical  
knowledge



Epistemic  
knowledge



# (finite) Markov Decision Process



trajectory

$$S_0, A_0, R_1, S_1, A_1, R_2, S_2, A_2, R_3, \dots$$

dynamics

$$p(s', r | s, a) \doteq \Pr\{S_t = s', R_t = r \mid S_{t-1} = s, A_{t-1} = a\}$$

$$\sum_{s' \in \mathcal{S}} \sum_{r \in \mathcal{R}} p(s', r | s, a) = 1, \text{ for all } s \in \mathcal{S}, a \in \mathcal{A}(s)$$

Perfect knowledge  
of the model

# (finite) Markov Decision Process

trajectory

$$S_0, A_0, R_1, S_1, A_1, R_2, S_2, A_2, R_3, \dots$$

dynamics

$$p(s', r | s, a) \doteq \Pr\{S_t = s', R_t = r \mid S_{t-1} = s, A_{t-1} = a\}$$

state-transition  
probability

$$p(s' | s, a) \doteq \Pr\{S_t = s' \mid S_{t-1} = s, A_{t-1} = a\} = \sum_{r \in \mathcal{R}} p(s', r | s, a)$$

expected reward (I)

$$r(s, a) \doteq \mathbb{E}[R_t \mid S_{t-1} = s, A_{t-1} = a] = \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} p(s', r | s, a)$$

expected reward (II)

$$r(s, a, s') \doteq \mathbb{E}[R_t \mid S_{t-1} = s, A_{t-1} = a, S_t = s'] = \sum_{r \in \mathcal{R}} r \frac{p(s', r | s, a)}{p(s' | s, a)}$$

# Reward signal



**Reward hypothesis:** that all of what we mean by goals and purposes can be well thought of as the maximization of the expected value of the cumulative sum of a received scalar signal (called reward).

Reward

$$R_{t+1}$$

Return

$$G_t \doteq R_{t+1} + R_{t+2} + R_{t+3} + \cdots + R_T$$

Discounted  
Return

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

# Reward signal



**Reward hypothesis:** that all of what we mean by goals and purposes can be well thought of as the maximization of the expected value of the cumulative sum of a received scalar signal (called reward).

Short-term view

Reward

$$R_{t+1}$$

Return

$$G_t \doteq R_{t+1} + R_{t+2} + R_{t+3} + \cdots + R_T$$

Long-term view

Discounted  
Return

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$



# Reward signal



**Reward hypothesis:** that all of what we mean by goals and purposes can be well thought of as the maximization of the expected value of the cumulative sum of a received scalar signal (called reward).

Short-term view

Reward

$$R_{t+1}$$

Return

$$G_t \doteq R_{t+1} + R_{t+2} + R_{t+3} + \cdots + R_T$$

Long-term view

Discounted  
Return

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$



$$G_t = R_{t+1} + \gamma G_{t+1}$$

**RECURSIVE  
DEFINITION**

# Policy

Policy

is a mapping from states to probabilities of selecting each possible action

$$\pi : \mathcal{S} \times \mathcal{A} \rightarrow [0, 1]$$

If we are at time  $t$ ,  $\pi(a|s)$  is the probability of having  $A_t = a \wedge S_t = s$

can be *deterministic*

# Value function

Value Function

is a function that quantify how good is to be on a state and follows a specific policy

$$v_{\pi} : \mathcal{S} \rightarrow \mathbb{R}$$

state-value  
function

$$v_{\pi}(s) \doteq \mathbb{E}_{\pi}[G_t \mid S_t = s] = \mathbb{E}_{\pi}\left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s\right], \text{ for all } s \in \mathcal{S}$$

action-value  
function

$$q_{\pi}(s, a) \doteq \mathbb{E}_{\pi}[G_t \mid S_t = s, A_t = a] = \mathbb{E}_{\pi}\left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s, A_t = a\right]$$

# Model of the environment



# Solving a RL problem

find a policy that achieves the maximum reward over the long run

optimal policy

$$\pi_* \succeq \pi \quad \forall \pi \in \text{policies}$$

# Solving a RL problem

find a policy that achieves the maximum reward over the long run

optimal policy

$$\pi_* \succeq \pi \quad \forall \pi \in \text{policies}$$

$$\pi' \succeq \pi \iff \forall s \in \mathcal{S}, v_{\pi'}(s) \geq v_{\pi}(s)$$

# Solving a RL problem

find a policy that achieves the maximum reward over the long run

optimal policy

$$\pi_* \succeq \pi \quad \forall \pi \in \text{policies}$$

$$\pi' \succeq \pi \iff \forall s \in \mathcal{S}, v_{\pi'}(s) \geq v_{\pi}(s)$$

optimal state-value  
function

$$v_*(s) \doteq \max_{\pi} v_{\pi}(s)$$

optimal action-value  
function

$$q_*(s, a) \doteq \max_{\pi} q_{\pi}(s, a)$$



# Solving a RL problem

find a policy that achieves the maximum reward over the long run

optimal policy

$$\pi_* \succeq \pi \quad \forall \pi \in \text{policies}$$

$$\pi' \succeq \pi \iff \forall s \in \mathcal{S}, v_{\pi'}(s) \geq v_{\pi}(s)$$

optimal state-value  
function

$$v_*(s) \doteq \max_{\pi} v_{\pi}(s)$$

optimal action-value  
function

$$q_*(s, a) \doteq \max_{\pi} q_{\pi}(s, a)$$

$$q_*(s, a) = \mathbb{E}[R_{t+1} + \gamma v_*(S_{t+1}) \mid S_t = s, A_t = a]$$



# Dynamic Programming


---

How to solve MDP problems


# Dynamic Programming

Mr. Richard Ernest Bellman

Algorithm paradigm useful to solve a specific class of problems that can be decomposed in sub-problems in recursive way



Bellman, 1950s




# Dynamic Programming

In the RL context

Collection of algorithms that can be used to compute optimal policies given a perfect model of the environment as a MDP.

**Key idea:** use value function to organize and structure the search of optimal policies



Bellman, 1950s

Consistency relation  
of state-value  
function

$$\begin{aligned} v_{\pi}(s) &\doteq \mathbb{E}_{\pi}[G_t \mid S_t = s] \\ &= \mathbb{E}_{\pi}[R_{t+1} + \gamma G_{t+1} \mid S_t = s] \\ &= \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) \left[ r + \gamma \mathbb{E}_{\pi}[G_{t+1} | S_{t+1} = s'] \right] \\ &= \sum_a \pi(a|s) \sum_{s', r} p(s', r | s, a) \left[ r + \gamma v_{\pi}(s') \right], \quad \text{for all } s \in \mathcal{S}. \end{aligned}$$

# Dynamic Programming

## Towards the Bellman Equation

Consistency relation  
of state-value  
function


$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$

# Dynamic Programming

## Towards the Bellman Equation

Consistency relation  
of state-value  
function

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$




# Dynamic Programming

## Towards the Bellman Equation

Consistency relation  
of state-value  
function

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$




# Dynamic Programming

## Towards the Bellman Equation

Consistency relation  
of state-value  
function

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$




# Dynamic Programming

## Towards the Bellman Equation

Consistency relation  
of state-value  
function

$$v_{\pi}(s) = \sum_a \pi(a|s) q_{\pi}(s, a)$$






# Dynamic Programming

## Towards the Bellman Equation

What about the optimal policy  
and the optimal state-value  
function?

Consistency relation  
of state-value  
function

$$v_{\pi}(s) = \sum_a \pi(a|s) q_{\pi}(s, a)$$



# Dynamic Programming


## Towards the Bellman Equation

Consistency relation  
of state-value  
function

$$v_{\pi}(s) = \sum_a \pi(a|s) q_{\pi}(s, a)$$

What about the optimal policy  
and the optimal state-value  
function?

It's an average




# Dynamic Programming

## Towards the Bellman Equation

Consistency relation  
of state-value  
function

$$v_{\pi}(s) = \sum_a \pi(a|s) q_{\pi}(s, a)$$



What about the optimal policy  
and the optimal state-value  
function?

It's an average


The optimal policy is a policy so  
it should satisfy the consistency  
relation

# Dynamic Programming

## Towards the Bellman Equation

Consistency relation  
of state-value  
function

$$v_*(s) = \sum_a \pi_*(a|s) q_*(s, a)$$



What about the optimal policy  
and the optimal state-value  
function?

It's an average


The optimal policy is a policy so  
it should satisfy the consistency  
relation

# Dynamic Programming

## Towards the Bellman Equation

Consistency relation  
of state-value  
function

$$v_*(s) = \sum_a \pi_*(a|s) q_*(s, a)$$



What about the optimal policy  
and the optimal state-value  
function?

It's an average

The optimal policy is a policy so  
it should satisfy the consistency  
relation


The optimal policy is *optimal*

# Dynamic Programming

## Towards the Bellman Equation

Consistency relation  
of state-value  
function

$$v_*(s) = \sum_a \pi_*(a|s) q_*(s, a)$$



What about the optimal policy  
and the optimal state-value  
function?

It's an average

The optimal policy is a policy so  
it should satisfy the consistency  
relation


The optimal policy is *optimal*

# Dynamic Programming

## Bellman Equation

Bellman equation

$$v_*(s) = \max_a q_*(s, a)$$




# Dynamic Programming

## Bellman Equation

Bellman equation

$$v_*(s) = \max_a q_*(s, a)$$


$$v_*(s) = \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_*(s')]$$






# Dynamic Programming

## Bellman Equation



# Dynamic Programming


How to find the optimal policy?



# Dynamic Programming

How to find the optimal policy?

Consistency relation of the  
state-value function




# Dynamic Programming

How to find the optimal policy?

Consistency relation of the  
state-value function

Bellman intuition




# Dynamic Programming

How to find the optimal policy?

Consistency relation of the  
state-value function

Bellman intuition




Policy Iteration

# Dynamic Programming

How to find the optimal policy?

Consistency relation of the  
state-value function

Bellman intuition



Policy Iteration

$$\pi_0 \xrightarrow{\text{E}} v_{\pi_0} \xrightarrow{\text{I}} \pi_1 \xrightarrow{\text{E}} v_{\pi_1} \xrightarrow{\text{I}} \pi_2 \xrightarrow{\text{E}} \dots \xrightarrow{\text{I}} \pi_* \xrightarrow{\text{E}} v_*$$

Does it converge? Yes

# Dynamic Programming

## Policy evaluation

$$\pi \xrightarrow{\text{Policy evaluation}} v_\pi$$

Consistency relation  
of state-value  
function

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_\pi(s')]$$

# Dynamic Programming

## Policy evaluation

$$\pi \xrightarrow{\text{Policy evaluation}} v_\pi$$

Consistency relation  
of state-value  
function

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_\pi(s')]$$

Iterative policy  
evaluation

$$v_{k+1}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_k(s')]$$



# Dynamic Programming

## Policy evaluation

$$\pi \xrightarrow{\text{Policy evaluation}} v_\pi$$

Consistency relation  
of state-value  
function


$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_\pi(s')]$$

Iterative policy  
evaluation

$$v_{k+1}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_k(s')]$$

2 ways of updating: in-place vs two  
arrays version

Faster, depends on ordering of  
update



# Policy Evaluation

## Algorithm

### Iterative Policy Evaluation, for estimating $V \approx v_\pi$

Input  $\pi$ , the policy to be evaluated

Algorithm parameter: a small threshold  $\theta > 0$  determining accuracy of estimation

Initialize  $V(s)$ , for all  $s \in \mathcal{S}^+$ , arbitrarily except that  $V(\text{terminal}) = 0$

Loop:

$\Delta \leftarrow 0$

Loop for each  $s \in \mathcal{S}$ :

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until  $\Delta < \theta$

# Policy Evaluation

## Algorithm

### Iterative Policy Evaluation, for estimating $V \approx v_\pi$

Input  $\pi$ , the policy to be evaluated

Algorithm parameter: a small threshold  $\theta > 0$  determining accuracy of estimation

Initialize  $V(s)$ , for all  $s \in \mathcal{S}^+$ , arbitrarily except that  $V(\text{terminal}) = 0$

Loop:

$\Delta \leftarrow 0$

Loop for each  $s \in \mathcal{S}$ :

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

consistency relation

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until  $\Delta < \theta$

# Policy Evaluation

## Algorithm

### Iterative Policy Evaluation, for estimating $V \approx v_\pi$

Input  $\pi$ , the policy to be evaluated

Algorithm parameter: a small threshold  $\theta > 0$  determining accuracy of estimation

Initialize  $V(s)$ , for all  $s \in \mathcal{S}^+$ , arbitrarily except that  $V(\text{terminal}) = 0$

Loop:

$\Delta \leftarrow 0$

Loop for each  $s \in \mathcal{S}$ :

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')]$

consistency relation

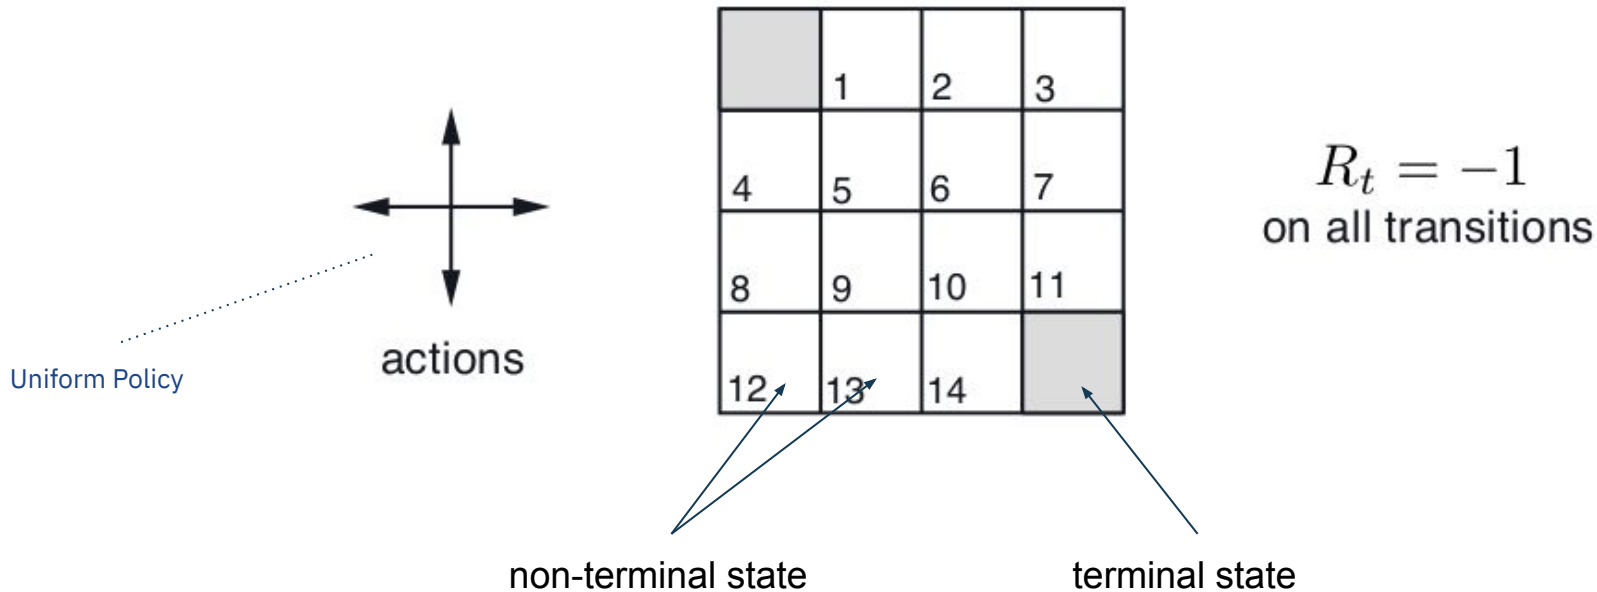
$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until  $\Delta < \theta$

Stability of state-value function


# Policy Evaluation

Example



# Policy Evaluation


Example - 1st iteration



$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$

# Policy Evaluation

Example - 1st iteration




$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [\underbrace{r}_{-1} + \gamma \underbrace{v_{\pi}(s')}_0]$$

# Policy Evaluation

Example - 1st iteration

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0




0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$



# Policy Evaluation

Example - 1st iteration



$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$

# Policy Evaluation

Example - 2nd iteration

← -1/3 +

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0




0.0	?		
?			
			?
		?	0.0

$$v_{\pi}(s) = \sum_a \frac{\pi(a|s)}{1/3} \sum_{s',r} \frac{p(s',r|s,a)}{1} \frac{r}{-1} + \gamma \frac{v_{\pi}(s')}{0}$$

# Policy Evaluation

Example - 2nd iteration



$$v_{\pi}(s) = \sum_a \frac{\pi(a|s)}{1/3} \sum_{s',r} \frac{p(s',r|s,a)}{1} \frac{[r + \gamma v_{\pi}(s')]}{-1}$$

# Policy Evaluation

Example - 2nd iteration

$\leftarrow$  -1/3 +  
 $\downarrow$  -2/3 +  
 $\rightarrow$  -2/3 +

-1.7

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0




0.0	?		
?			
			?
		?	0.0

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a)[r + \gamma v_{\pi}(s')]$$

# Policy Evaluation

Example - 2nd iteration



$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$

# Policy Evaluation

Example - 2nd iteration

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0



0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$

# Policy Evaluation

Example - until the end

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

1  
→

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

2  
→

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

3  
→

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0



0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

∞  
←

# Policy Evaluation

Example - until the end

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

1  
→

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

2  
→

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

3  
→

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0



0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

∞  
←



# Policy Evaluation

Example - until the end

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

1  
→

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

2  
→

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

3  
→

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0



0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

8  
←

# Policy Evaluation

Example - until the end

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

1  
→

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

2  
→


0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

3  
→

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0




0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0



# Policy Improvement

How to find better policies



0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

Policy improvement theorem


$$q_\pi(s, \pi'(s)) \geq v_\pi(s), \forall s \in \mathcal{S} \Rightarrow v_{\pi'}(s) \geq v_\pi(s), \forall s \in \mathcal{S}$$

Greedy policy approach

$$\begin{aligned}\pi'(s) &\doteq \arg \max_a q_\pi(s, a) \\ &= \arg \max_a \mathbb{E}[R_{t+1} + \gamma v_\pi(S_{t+1}) \mid S_t = s, A_t = a] \\ &= \arg \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_\pi(s')],\end{aligned}$$

# Policy Iteration


## Example



# Policy Iteration

## Example

propagation effect



policy convergence

# Policy Iteration

Example

propagation effect








0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0



policy convergence

# Policy Iteration

Policy Iteration (using iterative policy evaluation) for estimating  $\pi \approx \pi_*$

## 1. Initialization

$V(s) \in \mathbb{R}$  and  $\pi(s) \in \mathcal{A}(s)$  arbitrarily for all  $s \in \mathcal{S}$

## 2. Policy Evaluation

Loop:

$\Delta \leftarrow 0$

Loop for each  $s \in \mathcal{S}$ :

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_{s',r} p(s',r|s,\pi(s)) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until  $\Delta < \theta$  (a small positive number determining the accuracy of estimation)

## 3. Policy Improvement

*policy-stable*  $\leftarrow$  true

For each  $s \in \mathcal{S}$ :

*old-action*  $\leftarrow \pi(s)$

$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

If *old-action*  $\neq \pi(s)$ , then *policy-stable*  $\leftarrow$  false

If *policy-stable*, then stop and return  $V \approx v_*$  and  $\pi \approx \pi_*$ ; else go to 2



# Policy Iteration

Policy Iteration (using iterative policy evaluation) for estimating  $\pi \approx \pi_*$

## 1. Initialization

$V(s) \in \mathbb{R}$  and  $\pi(s) \in \mathcal{A}(s)$  arbitrarily for all  $s \in \mathcal{S}$

## 2. Policy Evaluation

Loop:

$\Delta \leftarrow 0$

Loop for each  $s \in \mathcal{S}$ :

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_{s',r} p(s',r|s,\pi(s)) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until  $\Delta < \theta$  (a small positive number determining the accuracy of estimation)

## 3. Policy Improvement

*policy-stable*  $\leftarrow$  true

For each  $s \in \mathcal{S}$ :

*old-action*  $\leftarrow \pi(s)$

$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

If *old-action*  $\neq \pi(s)$ , then *policy-stable*  $\leftarrow$  false

If *policy-stable*, then stop and return  $V \approx v_*$  and  $\pi \approx \pi_*$ ; else go to 2





# Value Iteration

Solving efficiently the Policy Iteration

## Value Iteration, for estimating $\pi \approx \pi_*$

Algorithm parameter: a small threshold  $\theta > 0$  determining accuracy of estimation

Initialize  $V(s)$ , for all  $s \in \mathcal{S}^+$ , arbitrarily except that  $V(\text{terminal}) = 0$


Loop:

```
|  $\Delta \leftarrow 0$   
| Loop for each  $s \in \mathcal{S}$ :  
|    $v \leftarrow V(s)$   
|    $V(s) \leftarrow \max_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$   
|    $\Delta \leftarrow \max(\Delta, |v - V(s)|)$   
until  $\Delta < \theta$ 
```

Output a deterministic policy,  $\pi \approx \pi_*$ , such that

$$\pi(s) = \arg \max_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$$

# Recap





# Thank you

[esteco.com](https://www.esteco.com)

